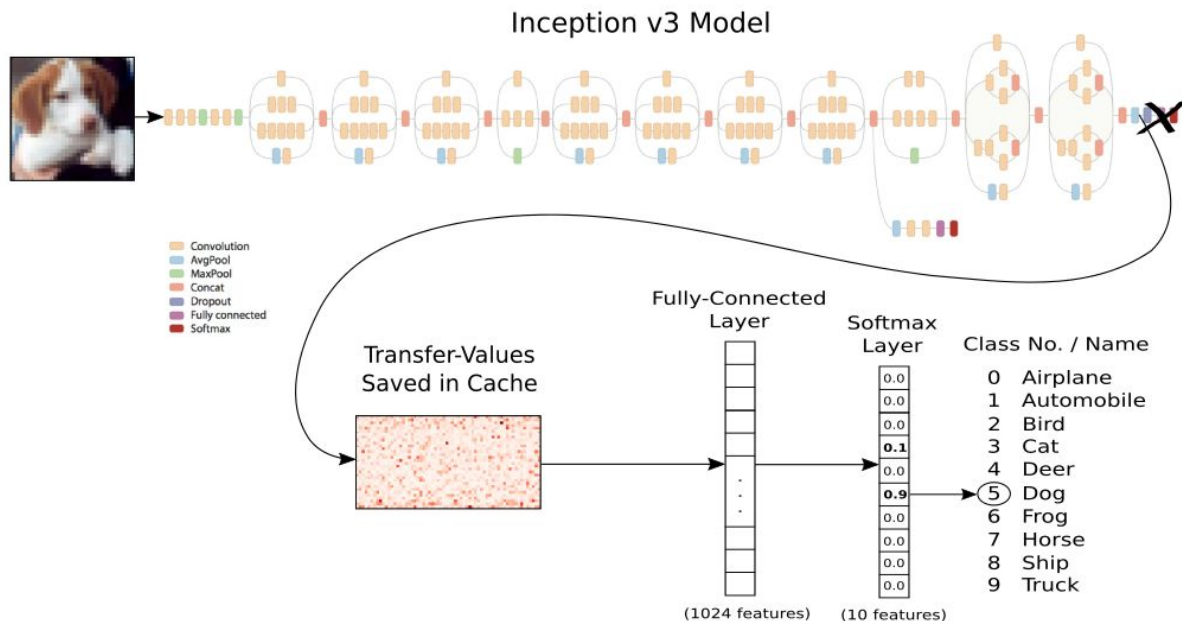


Understanding Black-box Predictions via Influence Functions

— Pang Wei Koh, Percy Liang —


As narrated by: Dany Haddad, Alan Gee

Big Picture



How can we explain where results of the model came from?

Influence Functions



Main Idea: quantify the influence of a training sample on the loss at a test sample. Not actually a black-box method.

Applications:

- Determine which training examples are most beneficial or harmful to test set performance
- Prioritizes which training set points to check for errors
- Generating adversarial *training* examples

Efficiency:

- No need to retrain the entire model

Assumptions

- Strictly convex and twice differentiable objective function
- What if these assumptions don't hold?
 - Can form a local strictly convex quadratic approximation

$$L(z, \theta) \approx L(z, \tilde{\theta}) + \nabla L(z, \tilde{\theta})^T (\theta - \tilde{\theta}) + \frac{1}{2} (\theta - \tilde{\theta})^T (H_{\tilde{\theta}} + \lambda I) (\theta - \tilde{\theta})$$

- For the Hinge loss case they propose to compute the Influence function through a smooth hinge loss term

$$\text{SmoothHinge}(s, t) = t \log(1 + \exp(\frac{1-s}{t}))$$

Influence Functions: Mathematical Formulation

$$z_i = (x_i, y_i) \quad \hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{n} \sum_i^n L(z_i, \theta)$$

Upweight the influence of z_j by some ϵ . The new model parameters are:

$$\hat{\theta}_{\epsilon, z} = \operatorname{argmin}_{\theta} \frac{1}{n} \sum_i^n L(z_i, \theta) + \epsilon L(z_j, \theta)$$

This has an influence on the parameters given by [Cook & Weisberg, 1982]:

$$\mathcal{I}_{\text{up, params}}(z) \stackrel{\text{def}}{=} \left. \frac{d\hat{\theta}_{\epsilon, z}}{d\epsilon} \right|_{\epsilon=0} = -H_{\hat{\theta}}^{-1} \nabla_{\theta} L(z, \hat{\theta}),$$

$$H_{\hat{\theta}} \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \nabla_{\theta}^2 L(z_i, \hat{\theta})$$



Influence on Test Loss

$$\mathcal{I}_{\text{up,loss}}(z, z_{\text{test}}) = -\nabla_{\theta} L(z_{\text{test}}, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_{\theta} L(z, \hat{\theta})$$

H_{θ}^{-1} tells us that if $\nabla_{\theta} L(z, \theta)$ points in a direction of little curvature, its influence on the loss at z_{test} will be higher

Influence Functions vs nearest neighbors

$$\mathcal{I}_{\text{up,loss}}(z, z_{\text{test}}) = -y_{\text{test}}y \cdot \sigma(-y_{\text{test}}\theta^\top x_{\text{test}}) \cdot \sigma(-y\theta^\top x) \cdot x_{\text{test}}^\top H_{\hat{\theta}}^{-1}x.$$



Efficiently Computing Influence

Main Idea: Can compute $s_{\text{test}} \stackrel{\text{def}}{=} H_{\hat{\theta}}^{-1} \nabla_{\theta} L(z_{\text{test}}, \hat{\theta})$ without explicitly inverting H_{θ}

Conjugate gradient method: Compute s_{test} by framing it as a solution to a system of equations


$$H_{\hat{\theta}} s_{\text{test}} = \nabla_{\theta} L(z_{\text{test}}, \hat{\theta})$$

Stochastic estimation [Agarwal et al. (2016)]: Sample t training points uniformly and recursively compute an unbiased estimate of s_{test} . Empirically show this is faster than CG

$$\hat{s}_{\text{test},j} = \nabla_{\theta} L(z_{\text{test}}, \hat{\theta}) + (I - \nabla_{\theta}^2 L(z_j, \hat{\theta})) \hat{s}_{\text{test},j-1}$$

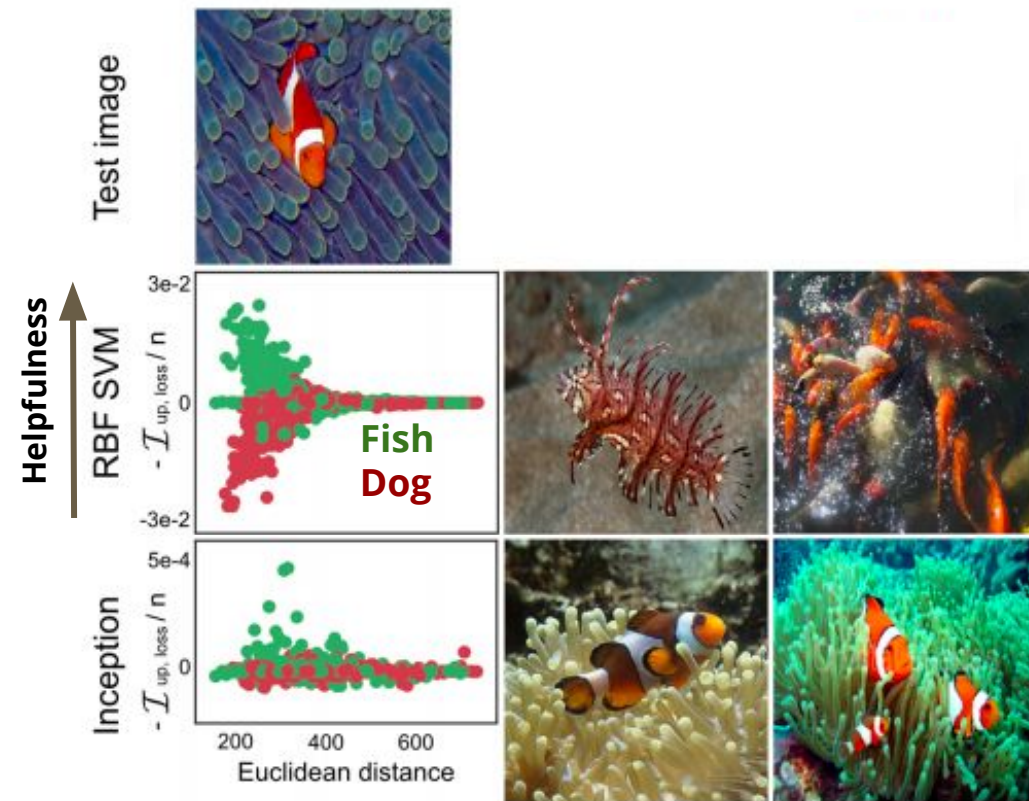
Comparison to Leave-one-out Training

Well correlated with actual difference in loss even for more complex models (such as a CNN)



Applications

Understanding Model Behavior



Influence functions can show which training data can help make a correct prediction on a test image

For example: the dog (pictured) was most helpful in determining that the test image was a fish

Generating examples for Adversarial Attacks

Small, and in direction of low variance

Label: Fish






A small perturbation to one training example:

$+ \epsilon \cdot$

Label: Fish

Ambiguous examples are easy to attack

Can change multiple **test** predictions:

				
Orig (confidence): Dog (97%)	Dog (98%)	Dog (98%)	Dog (99%)	Dog (98%)
New (confidence): Fish (97%)	Fish (93%)	Fish (87%)	Fish (63%)	Fish (52%)

Attacking one image flipped 57% of the test predictions, 2 images flipped 77%

Finding training points that have most influence

Domain Mismatch


Training data distribution doesn't match test data distribution

Example: data from different hospitals serve different populations

- Use I.F. to track training data that drive incorrect classifications
- Example: identifying offenders in an Imbalanced data set (e.g. children are re-admitted with high % so classifier predicts most re-admits are children)

Mislabeled Examples

Focusing the attention of experts on training data that may be "corrupt": E.g. can we find the randomly flipped labels (~10%) in training data for spam filtering



Code and Data Repository

Dockerfile (full tensorflow environment): <https://hub.docker.com/r/pangwei/tf1.1/>

- Numpy/Scipy/Scikit-learn/Pandas
 - Tensorflow (tested on v1.1.0)
 - Keras (tested on v2.0.4)
- Spacy (tested on v1.8.2)
h5py (tested on v2.7.0)

Datasets

We use the MNIST, Dogfish, Enron spam, and UCI diabetes hospital readmission datasets. For attribution, please see the paper.

uuid[0:8]	name	summary	data_size	state	description
0xb2d85c	mnist	[uploaded]	11.1m	ready	MNIST
0x550cd3	dogfish	[uploaded]	741m	ready	Dogfish
0x19a9fd	spam	[uploaded]	7.2m	ready	Spam (enron1) data
0xfabc11	hospital	[uploaded]	18.3m	ready	Hospital readmission data

GitHub: <http://bit.ly/gt-influence>

Codalab: <http://bit.ly/cl-influence>

Discussion

Main contribution: Use of Influence functions to perturb model parameters and training data to realize the impact on the test data

Influence functions give us explainability but not interpretability

- Helps with model and dataset debugging after model parameters have been generated

Influence functions measure the effect of local changes: what happens when we upweight a point with small perturbations?

- Lack of global influence

Can be applied to models with non-differentiable loss functions by replacing with smooth loss functions

As developed by the authors, which of the following terms is NOT a component that effects $I_{up,loss}$, the influence of upweighting z on the loss at a test point z_{test} ?

- a) training loss
- b) weighted covariance matrix
- c) scaled Euclidean distance
- d) nearest neighbor in Euclidean space

Calculating the influence function can be computationally challenging. Which method did the authors contribute to to improve computational efficiency for large data sets?

- a) conjugate gradients
- b) inversion of the Hessian of the empirical risk function
- c) second-order stochastic estimation
- d) leave-one-out retraining

What was NOT a model used by the authors to demonstrate the significance of influence functions?

- a) Random forest
- b) Logistic regression
- c) Inception v3 neural network
- d) SVM with RBF kernels

What was NOT a use-case offered by the authors of the application of influence functions?

- a) Identifying models that heavily rely on few training examples
- b) Removing training points to improve classification predictions
- c) Identifying training examples responsible for misclassification errors
- d) Fixing mislabeled examples

Identifying data points for Adversarial Attacks

Use influence functions to generate small perturbations on training set that change prediction outcomes on test images

$$z = (x, y) \mapsto z_{\delta} = (x + \delta, y)$$

$\mathcal{I}_{\text{pert,loss}}(z, z_{\text{test}})^{\top}$: shows how we can perturb a training point z by δ to have the maximal increase in the loss on z_{test}

$$\mathcal{I}_{\text{pert,loss}}(z, z_{\text{test}})^{\top} = -\nabla_{\theta} L(z_{\text{test}}, \hat{\theta})^{\top} H_{\hat{\theta}}^{-1} \nabla_x \nabla_{\theta} L(z, \hat{\theta}).$$