

# Towards a Fair Marketplace

R. Mehrotra<sup>1</sup>, J. McInerney<sup>1</sup>, H. Bouchard<sup>1</sup>, M. Lalmas<sup>1</sup>, F. Diaz<sup>2</sup>

<sup>1</sup>Spotify Research, <sup>2</sup>Microsoft Research

As narrated by: Alan Gee  
and Dany Haddad

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

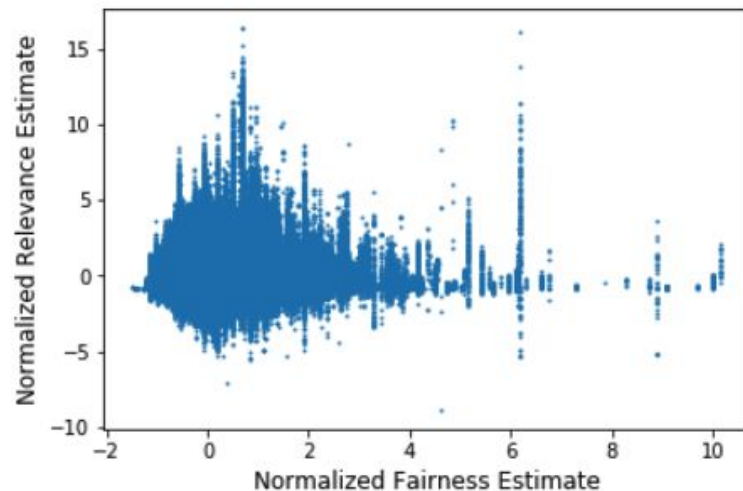
# Paper Roadmap

- Recommendation systems as 2-sided marketplace (suppliers, consumers)
- Fairness towards the suppliers to avoid *superstar economics*
  - *S.E. : boosts popular choices, impedes new or less popular choices*
- Propose a *group fairness measure  $\psi$*  , and other metrics
  - Take advantage of some users being more open accepting of lower relevance results
- Evaluate unbiased user satisfaction using counterfactual estimation instead of A/B testing

# Problem Setup

- Music recommendation and streaming service
- Goal is to maximize user satisfaction
  - Maximize the number of songs a user listens to

**Key question: How can we be fair in recommending less popular suppliers to consumers but maintain good relevance?**



# Learning Relevance (Music) Recommendation

- Train skip-gram vector for each song and user based on historical listening logs
- Relevance of a track (set of tracks) to a user is given by cosine-similarity (average cosine-similarity)


$$\sum_{u \in \mathcal{U}} \sum_{t \in \mathcal{P}_u} \log p(t|u)$$

$$p(t|u) = \frac{u^T t}{\sum_{t' \in \mathcal{T}} u^T t'}$$


# Fairness Measure

- Artists are binned into groups based on popularity
- Measure resembles a set-cover penalty, encouraging songs from artists from each popularity group to be represented

$$\psi(s) = \sum_i^K \sqrt{|G_i|}$$



$$\psi(s_1) = \sqrt{2} + \sqrt{2} = 2.8$$



$$\psi(s_2) = \sqrt{3} + \sqrt{1} = 2.7$$

$$\psi(s_1) = \sqrt{2} + \sqrt{2} = 2.8 > \psi(s_2) = \sqrt{3} + \sqrt{1} = 2.7$$

# Relevance vs Fairness trade-off

- Optimizing Relevance (personalization of recommendations for customers)
  - Learned embeddings from historical logs

**Optimal set for user**

$$s_u^* = \operatorname{argmax}_{s \in S} \phi(u, s)$$

**Relevance function**



- Optimizing Fairness (showing content across popularity spectrum for suppliers)

$$s_u^* = \operatorname{argmax}_{s \in S_u} \psi(s)$$

**Sets pertinent to user**



**Fairness function**



# Group Level: Combining Relevance and Fairness

- Weighted Mixture ( $\phi$ : relevance,  $\psi$ : fairness)

$$s_u^* = \operatorname{argmax}_{s \in S_u} ((1 - \beta) \phi(u, s) + \beta \psi(s)) \quad \beta \in [0, 1]$$

- Probabilistic Mixture (random draw  $p$ )

$$s_u^* = \begin{cases} \operatorname{argmax}_{s \in S_u} \psi(s) & \text{if } p < \beta \quad p \in [0, 1] \\ \operatorname{argmax}_{s \in S} \phi(u, s) & \text{otherwise} \end{cases}$$

Choose a random  $p$  and see where it compares to tolerance  $\beta$

# Group Level: Combining Relevance and Fairness

- Guaranteeing Minimum Relevance with Fairness

$$\begin{aligned} s_u^* &= \operatorname{argmax}_{s \in S_u} \psi(s) \\ \text{s.t. } &\phi(s, u) \geq \beta \end{aligned}$$

Constraint ensures the minimum relevance of recommendation


- Allows for recommendation sets with less relevance (due to fairness)



# Personalized Recommendation

- Defined a user's affinity for fair content as\*:

$$\xi_u = \frac{1}{|N_{>\alpha_1}|} \underbrace{\sum_{\phi(u,s) > \alpha_1} \zeta(u,s)}_{\text{Historic User Satisfaction with relevance content}} - \frac{1}{|N_{>\alpha_2}|} \underbrace{\sum_{\psi(u,s) > \alpha_2} \zeta(u,s)}_{\text{Historic User Satisfaction with fairness content}}$$



Score can then be used to understand a user's recommendation policy

\*believed there is a typo (fixed here)

# Adaptive Recommender: Managing Trade-off

- Adaptive - I (strict bifurcation of affinity scores)

$$s_u^* = \begin{cases} \operatorname{argmax}_{s \in S_u} \psi(s) & \text{if } \xi_u \geq 0 \\ \operatorname{argmax}_{s \in S} \phi(u, s) & \text{if } \xi_u < 0 \end{cases}$$

**Fairness optimized when scores are positive**

**Relevance optimized when scores are negative**

- Adaptive - II

$$\hat{\xi}_u = \frac{\xi_u - \mu_{\xi_u}}{\sigma_{\xi_u}}$$

**Normalized score across all users**

$$s_u^* = \operatorname{argmax}_{s \in S_u} ((1 - \hat{\xi}_u) \phi(u, s) + \hat{\xi}_u \psi(s))$$

**Reweight recommendation**

# Contextual Bandit Formulation

- A policy  $\pi$  gives a distribution over actions
- An action corresponds to recommending a specific set of songs
- The resulting reward is the satisfaction of a user (how many tracks they listened to)



# Contextual Bandit Formulation

- Value of a policy  $\pi$  where:
  - Context,  $x$ , is drawn from the distribution over states,  $D$ 
    - User and song features, time of day etc
  - Action,  $a$ , is chosen from distribution given by the policy,  $\pi$
  - Reward,  $r$ , received is based on the context and action chosen


$$\begin{aligned} V(\pi) &= E_{(x,r) \sim D}[r_{\pi(x)}] \\ &= E_{x \sim D} E_{a \sim \pi(\cdot|x)} E_{r \sim D(\cdot|x,a)}[r_{\pi(x)}] \end{aligned}$$

# Offline evaluation of recommendations

- Allows us to avoid A/B testing!
- Collect randomized data
  - Assign playlists uniformly at random for each user
    - Unclear how the subset is chosen or how it affects the results
  - Record the user's satisfaction (number of tracks listened to in playlist)

# Offline evaluation of recommendations

$$\hat{V}_{offline} = \sum_{\forall(x,a,r_a,p_a)} \frac{r_a I\{\pi(x) = a\}}{p_a}$$



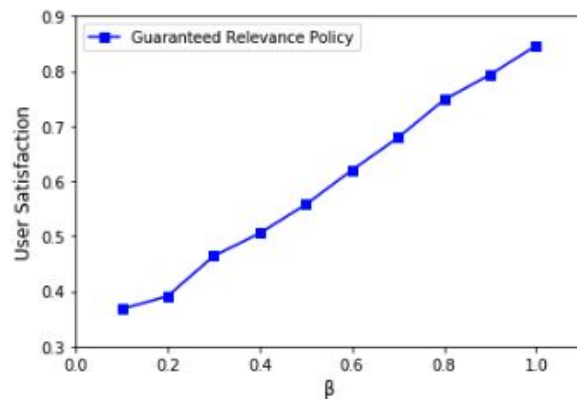
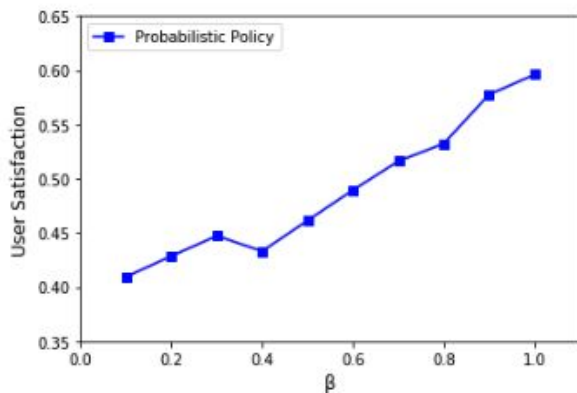
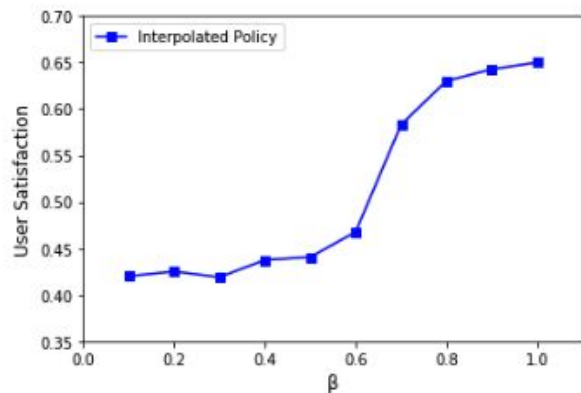
**Figure 2: Online vs. offline target metric values. Each point corresponds to one of the 7 days in the data collection period.**

# Tricks/Subtleties

- 400k users' rewards for 5k playlists from 50k artists
- Store random seed at each point to avoid issues from pseudo-random number generation

$$\mathbf{Var} \left[ \hat{V}_{\text{offline}}(\pi) \right] = \mathbf{E}_{(x, \tilde{r}) \sim D} \left[ r_{\pi(x)}^2 \left( \frac{1}{p_{\pi(x)}} - 1 \right) \right]$$

# Relevance vs Fairness Trade-off



- Favoring fairness over relevance decreases user satisfaction
  - Is this an acceptable trade-off?
- Guaranteed Relevance Policy seems to have the widest range of user satisfaction scores



# Performance of Recommendation Policies

Recommendation Policy	$\beta$	% Loss in Fairness	% Loss in Relevance	% Gain in Satisfaction
Only Fairness	N/A	0	57.7	-35.3
Only Relevant	N/A	69.1	0	0
Interpolated	0.5	3.32	48.7	-32.2
	0.7	42.7	9.8	-10.2
	0.9	64.7	0.06	-1.1
probPolicy	0.5	16.3	44.8	-29.0
	0.7	30.8	32.7	-20.6
	0.9	48.2	17.6	-11.1
GuaranteedR	0.5	37.7	19.6	-14.2
	0.7	51.7	7.8	4.4
	0.9	63.9	0.59	22.1
Adaptive - I	N/A	17	20.2	9.0
Adaptive - II	N/A	15	21.2	12.1

**Table 3: Comparing loss in fairness & relevance, with gains in satisfaction for different recommendation policies.**

# Summary

- Suggest approach to overcome *superstar economics*
- Give a heuristic to compute fairness towards a group of suppliers
- Provided a counterfactual estimation framework to estimate their impact on consumer satisfaction (offline)
- Show that some users are more accepting of “fair content” than others

# Fairness Measure: Alternative

- A more natural notion might be to assign inverse-popularity scores instead
  - Otherwise, you might just always recommend songs from the most popular artist in each group

